

MP4_P1

November 23, 2020

1 Generative Adversarial Networks

For this part of the assignment you implement two different types of generative adversarial networks. We will train the networks on the Celeb A dataset which is a large set of celebrity face images.

```
[3]: import torch
from torch.utils.data import DataLoader
from torchvision import transforms
from torchvision.datasets import ImageFolder

import numpy as np

import matplotlib.pyplot as plt
import matplotlib.gridspec as gridspec

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

%load_ext autoreload
%autoreload 2
```

```
[4]: from gan.train import train
```

```
[5]: device = torch.device("cuda:0" if torch.cuda.is_available() else "gpu")
```

2 GAN loss functions

In this assignment you will implement two different types of GAN cost functions. You will first implement the loss from the [original GAN paper](#). You will also implement the loss from [LS-GAN](#).

2.0.1 GAN loss

TODO: Implement the `discriminator_loss` and `generator_loss` functions in `gan/losses.py`.

The generator loss is given by:

$$\ell_G = -\mathbb{E}_{z \sim p(z)} [\log D(G(z))]$$

and the discriminator loss is:

$$\ell_D = -\mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] - \mathbb{E}_{z \sim p(z)} [\log (1 - D(G(z)))]$$

Note that these are negated from the equations presented earlier as we will be *minimizing* these losses.

HINTS: You should use the `torch.nn.functional.binary_cross_entropy_with_logits` function to compute the binary cross entropy loss since it is more numerically stable than using a softmax followed by BCE loss. The BCE loss is needed to compute the log probability of the true label given the logits output from the discriminator. Given a score $s \in \mathbb{R}$ and a label $y \in \{0, 1\}$, the binary cross entropy loss is

$$bce(s, y) = -y * \log(s) - (1 - y) * \log(1 - s)$$

Instead of computing the expectation of $\log D(G(z))$, $\log D(x)$ and $\log (1 - D(G(z)))$, we will be averaging over elements of the minibatch, so make sure to combine the loss by averaging instead of summing.

[4] : `from gan.losses import discriminator_loss, generator_loss`

2.0.2 Least Squares GAN loss

TODO: Implement the `ls_discriminator_loss` and `ls_generator_loss` functions in `gan/losses.py`.

We'll now look at [Least Squares GAN](#), a newer, more stable alternative to the original GAN loss function. For this part, all we have to do is change the loss function and retrain the model. We'll implement equation (9) in the paper, with the generator loss:

$$\ell_G = \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)) - 1)^2]$$

and the discriminator loss:

$$\ell_D = \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} [(D(x) - 1)^2] + \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)))^2]$$

HINTS: Instead of computing the expectation, we will be averaging over elements of the minibatch, so make sure to combine the loss by averaging instead of summing. When plugging in for $D(x)$ and $D(G(z))$ use the direct output from the discriminator (`scores_real` and `scores_fake`).

[5] : `from gan.losses import ls_discriminator_loss, ls_generator_loss`

3 GAN model architecture

TODO: Implement the Discriminator and Generator networks in `gan/models.py`.

We recommend the following architectures which are inspired by [DCGAN](#):

Discriminator:

- convolutional layer with in_channels=3, out_channels=128, kernel=4, stride=2
- convolutional layer with in_channels=128, out_channels=256, kernel=4, stride=2
- batch norm
- convolutional layer with in_channels=256, out_channels=512, kernel=4, stride=2
- batch norm
- convolutional layer with in_channels=512, out_channels=1024, kernel=4, stride=2
- batch norm
- convolutional layer with in_channels=1024, out_channels=1, kernel=4, stride=1

Use padding = 1 (not 0) for all the convolutional layers.

Instead of Relu we LeakyReLu throughout the discriminator (we use a negative slope value of 0.2). You can use simply use relu as well.

The output of your discriminator should be a single value score corresponding to each input sample. See `torch.nn.LeakyReLU`.

Generator:

Note: In the generator, you will need to use transposed convolution (sometimes known as fractionally-strided convolution or deconvolution). This function is implemented in pytorch as `torch.nn.ConvTranspose2d`.

- transpose convolution with in_channels=NOISE_DIM, out_channels=1024, kernel=4, stride=1
- batch norm
- transpose convolution with in_channels=1024, out_channels=512, kernel=4, stride=2
- batch norm
- transpose convolution with in_channels=512, out_channels=256, kernel=4, stride=2
- batch norm
- transpose convolution with in_channels=256, out_channels=128, kernel=4, stride=2
- batch norm
- transpose convolution with in_channels=128, out_channels=3, kernel=4, stride=2

The output of the final layer of the generator network should have a `tanh` nonlinearity to output values between -1 and 1. The output should be a 3x64x64 tensor for each sample (equal dimensions to the images from the dataset).

```
[6]: from gan.models import Discriminator, Generator
```

4 Data loading: Celeb A Dataset

The CelebA images we provide have been filtered to obtain only images with clear faces and have been cropped and downsampled to 128x128 resolution.

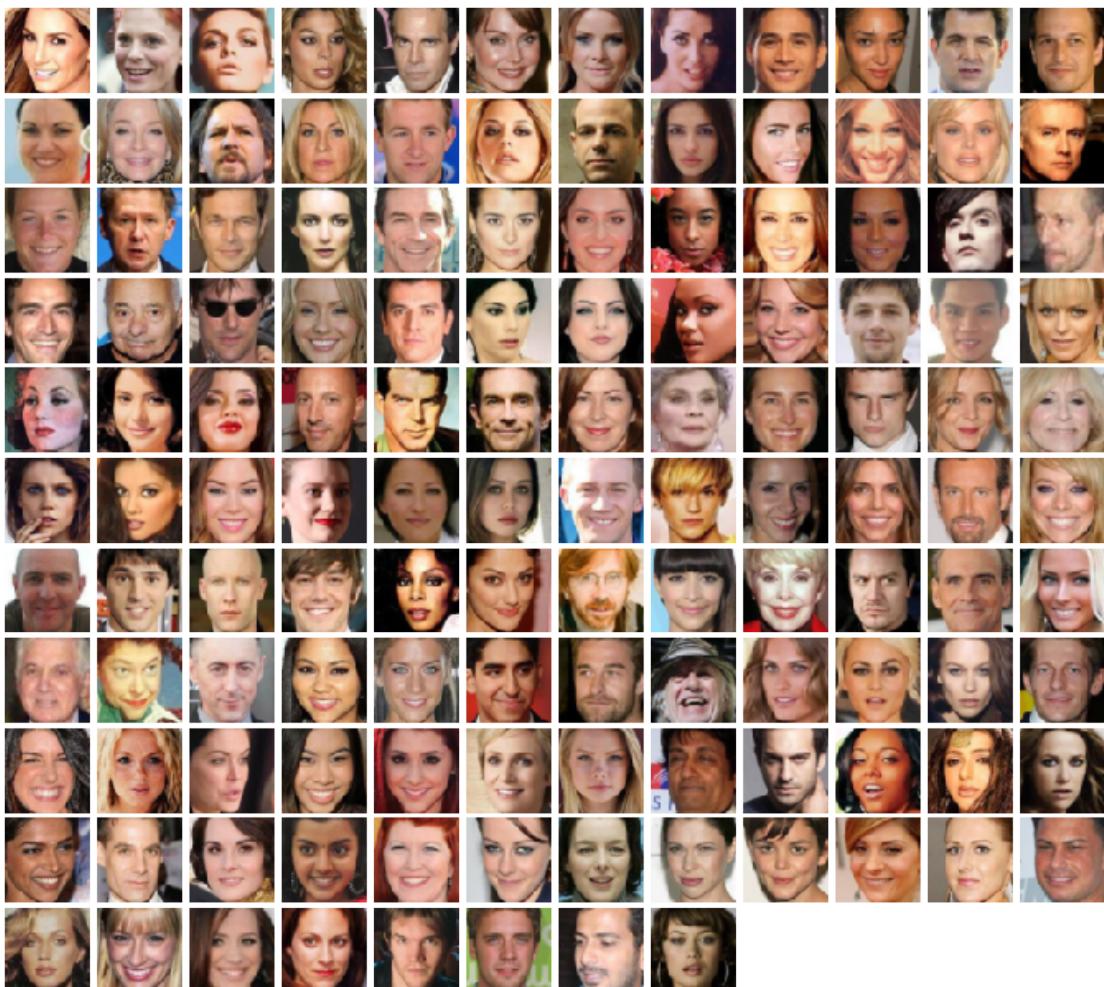
```
[6]: batch_size = 128
scale_size = 64 # We resize the images to 64x64 for training

celeba_root = 'celeba_data'
```

```
[7]: celeba_train = ImageFolder(root=celeba_root, transform=transforms.Compose([
    transforms.Resize(scale_size),
    transforms.ToTensor(),
]))  
  
celeba_loader_train = DataLoader(celeba_train, batch_size=batch_size, drop_last=True)
```

4.0.1 Visualize dataset

```
[9]: from gan.utils import show_images  
imgs = celeba_loader_train.__iter__().next()[0].numpy().squeeze()  
show_images(imgs, color=True)
```



5 Training

TODO: Fill in the training loop in `gan/train.py`.

```
[14]: NOISE_DIM = 100  
NUM_EPOCHS = 20  
learning_rate = 0.0002
```

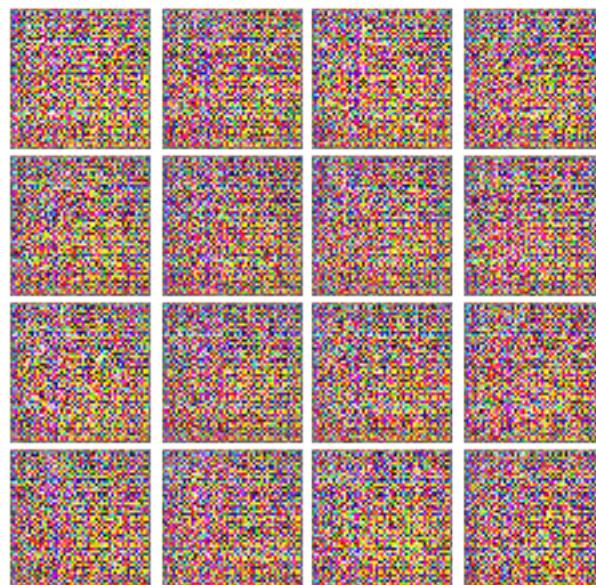
5.0.1 Train GAN

```
[32]: D = Discriminator().to(device)  
G = Generator(noise_dim=NOISE_DIM).to(device)
```

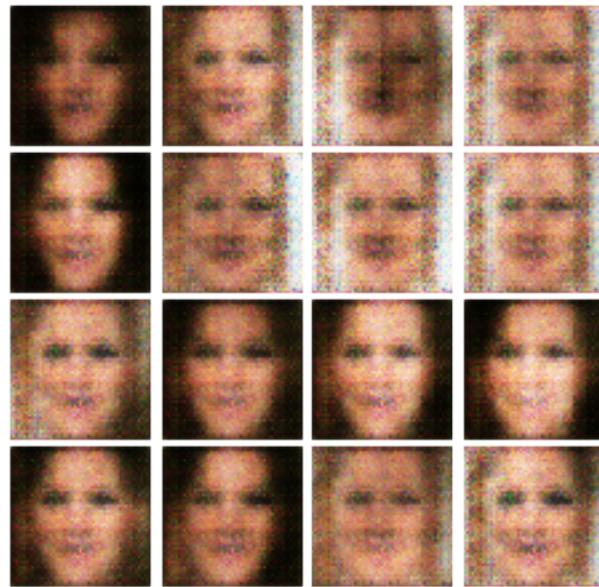
```
[33]: D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))  
G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))
```

```
[34]: # original gan  
train(D, G, D_optimizer, G_optimizer, discriminator_loss,  
      generator_loss, num_epochs=NUM_EPOCHS, show_every=200,  
      train_loader=celeba_loader_train, device=device)
```

EPOCH: 1
Iter: 0, D: 1.313, G:8.171



Iter: 200, D: 1.727, G:1.498



Iter: 400, D: 0.3473, G:3.814



Iter: 600, D: 0.2509, G:2.273

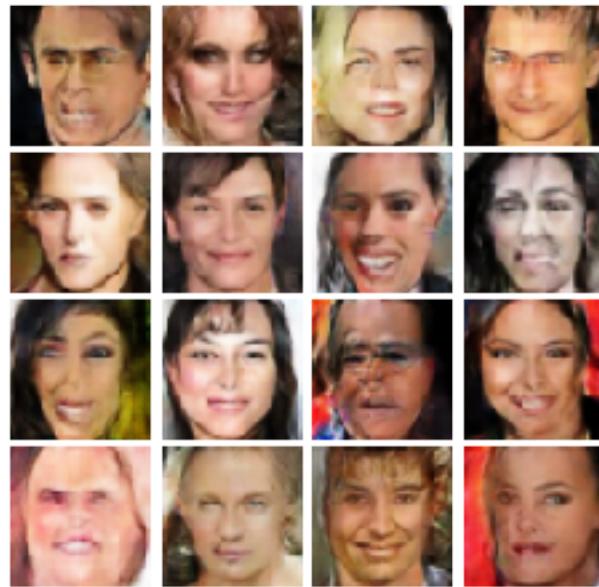


Iter: 800, D: 0.4293, G:2.302

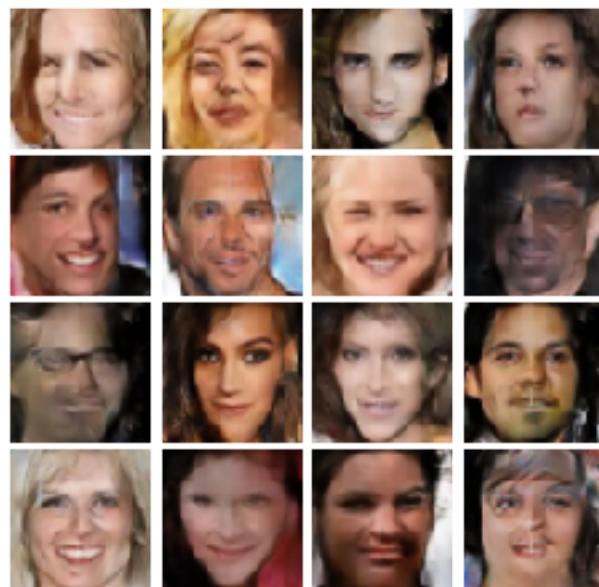


EPOCH: 2

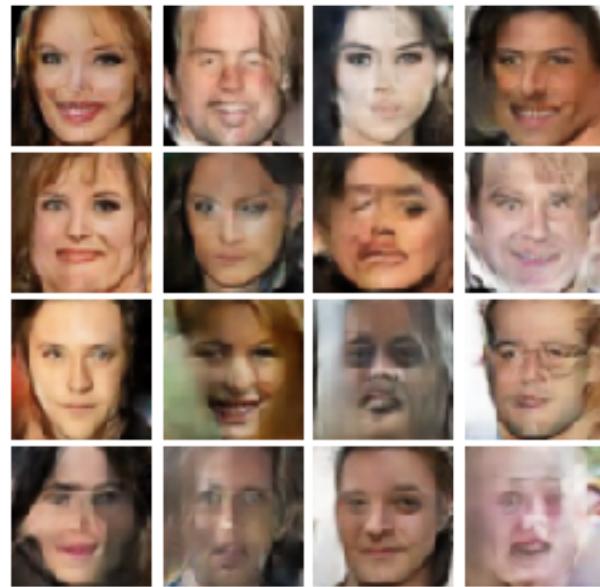
Iter: 1000, D: 0.3184, G:2.422



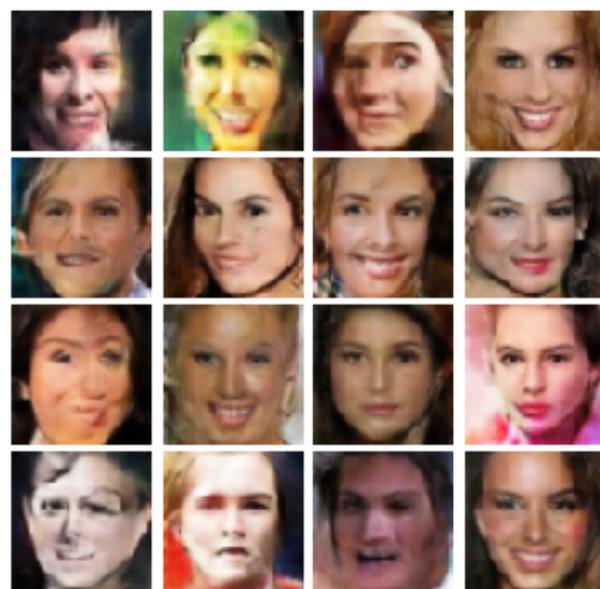
Iter: 19200, D: 0.03782, G:5.286



Iter: 19400, D: 0.1779, G:3.481



Iter: 19600, D: 0.1112, G:4.71



5.0.2 Train LS-GAN

```
[68]: D = Discriminator().to(device)
G = Generator(noise_dim=NOISE_DIM).to(device)

[69]: D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))
G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))

[70]: # ls-gan
train(D, G, D_optimizer, G_optimizer, ls_discriminator_loss,
      ls_generator_loss, num_epochs=NUM_EPOCHS, show_every=200,
      train_loader=celeba_loader_train, device=device)
```

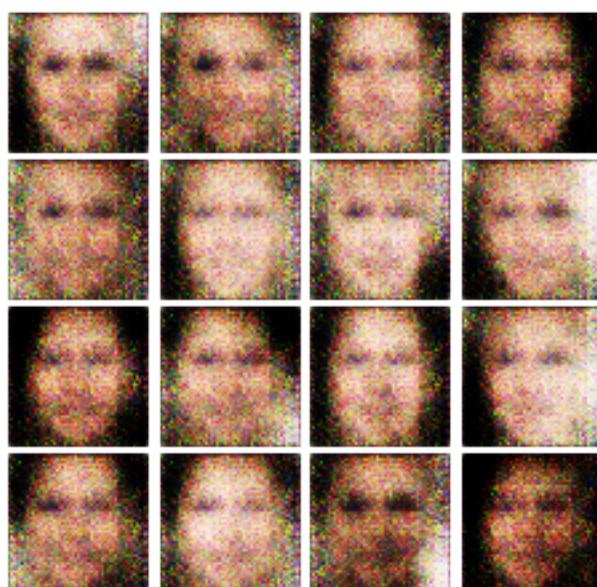
EPOCH: 1
Iter: 0, D: 1.753, G:40.06



Iter: 200, D: 0.05987, G:0.5976



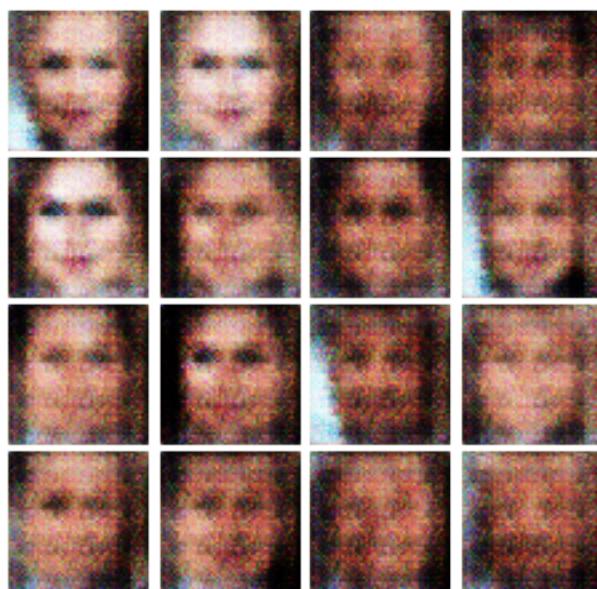
Iter: 400, D: 0.1178, G:0.1716



Iter: 600, D: 0.1241, G:0.2174

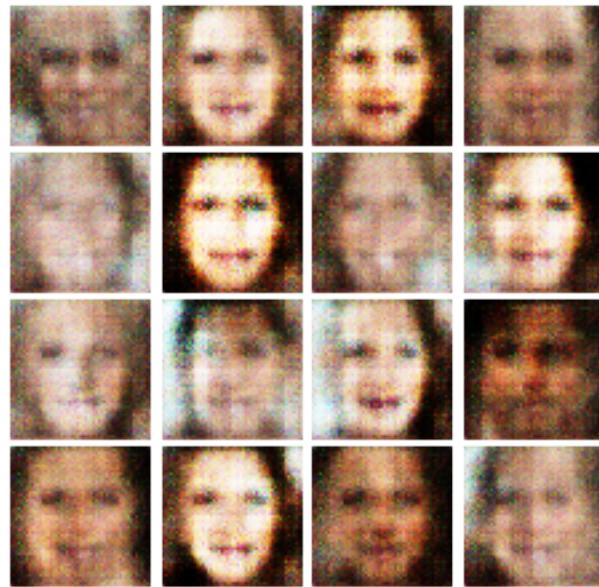


Iter: 800, D: 0.2117, G:0.4707

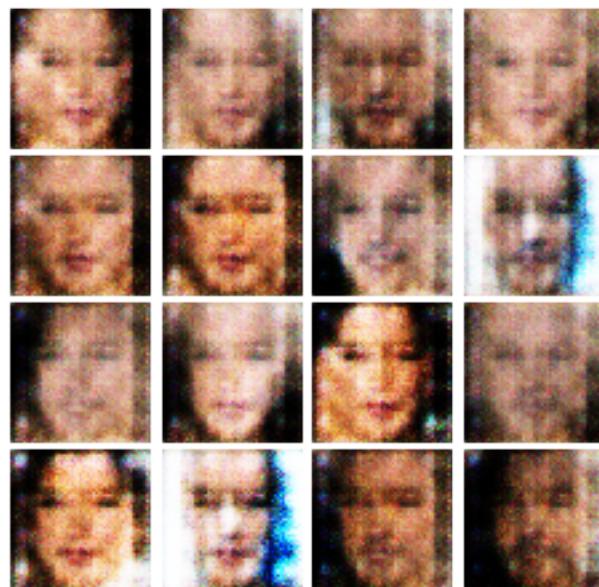


EPOCH: 2

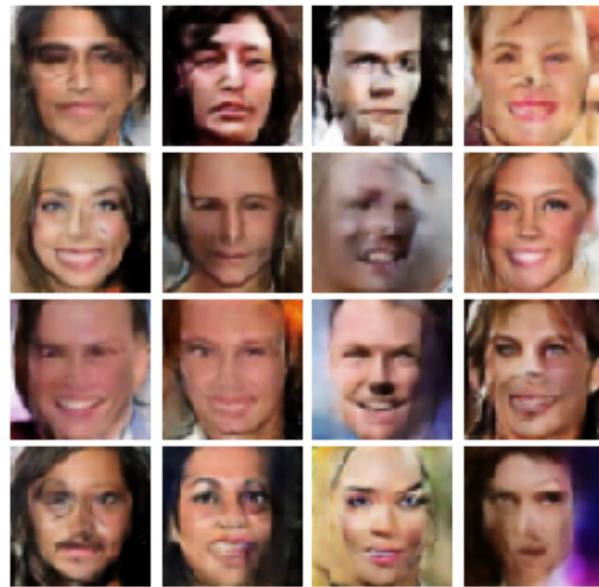
Iter: 1000, D: 0.1414, G:0.1573



Iter: 1200, D: 0.1481, G:0.1931



Iter: 1400, D: 0.2092, G:0.2114



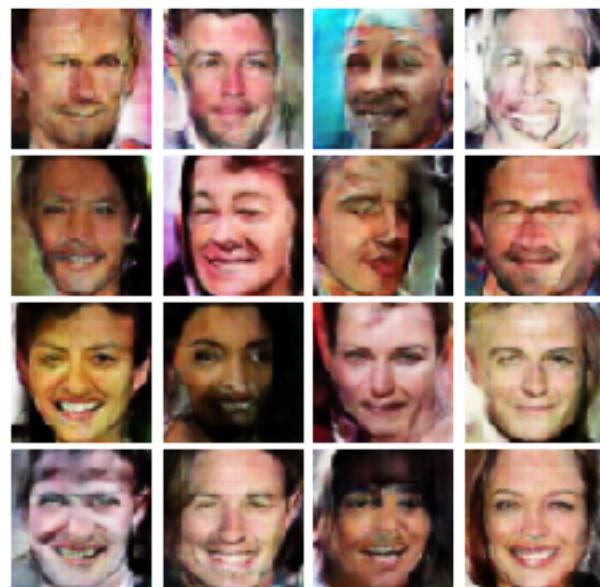
Iter: 19200, D: 0.04579, G: 0.5586



Iter: 19400, D: 0.03296, G: 0.4202



Iter: 19600, D: 0.05691, G: 0.6381



6 GAN Without Spectral Normalization

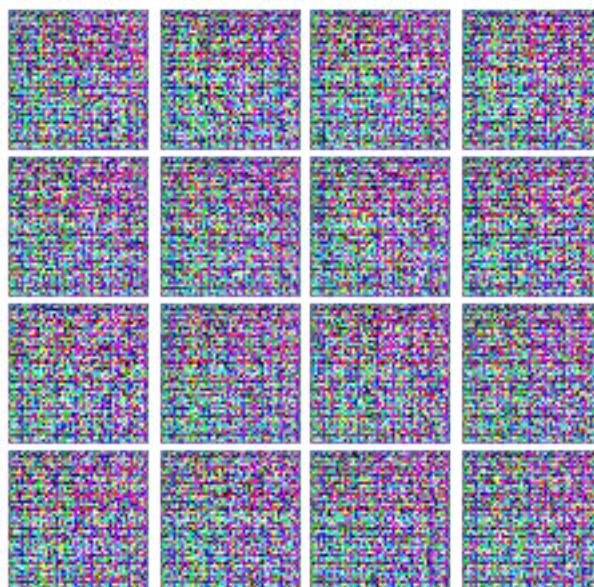
```
[72]: NUM_EPOCHS = 10

D = Discriminator().to(device)
G = Generator(noise_dim=NOISE_DIM).to(device)

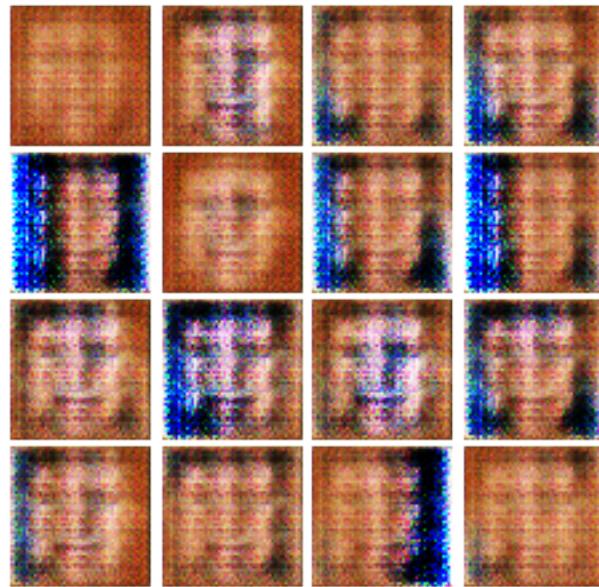
D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))
G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))

train(D, G, D_optimizer, G_optimizer, discriminator_loss,
      generator_loss, num_epochs=NUM_EPOCHS, show_every=200,
      train_loader=celeba_loader_train, device=device)
```

EPOCH: 1
Iter: 0, D: 1.371, G:3.29



Iter: 200, D: 0.4793, G:3.061



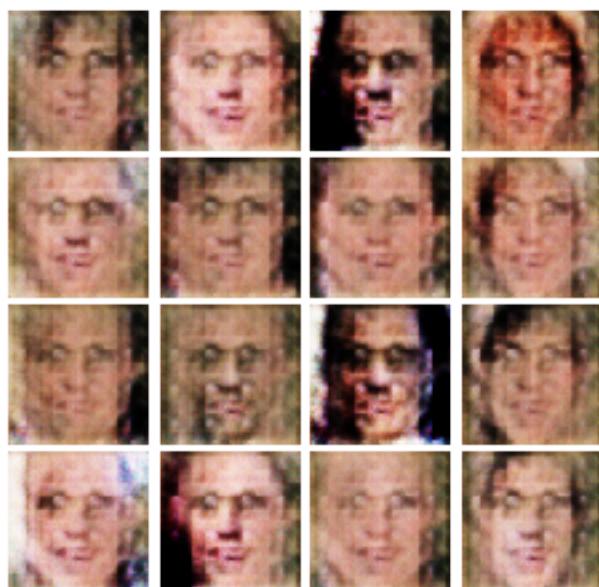
Iter: 400, D: 0.5463, G:2.378



Iter: 600, D: 1.787, G:4.888

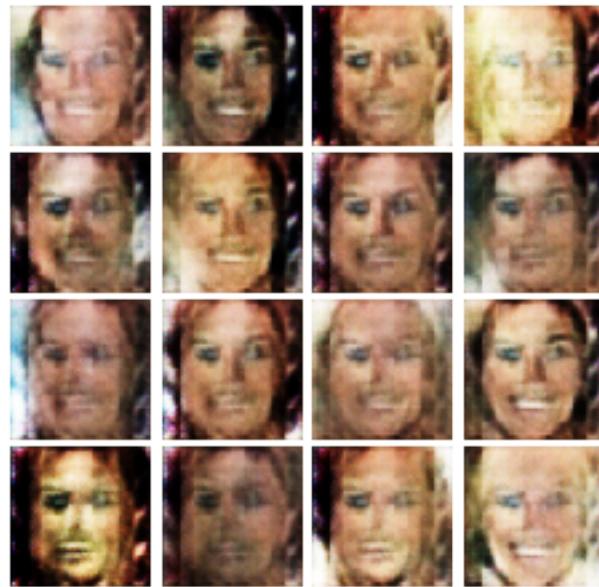


Iter: 800, D: 0.2592, G:2.287

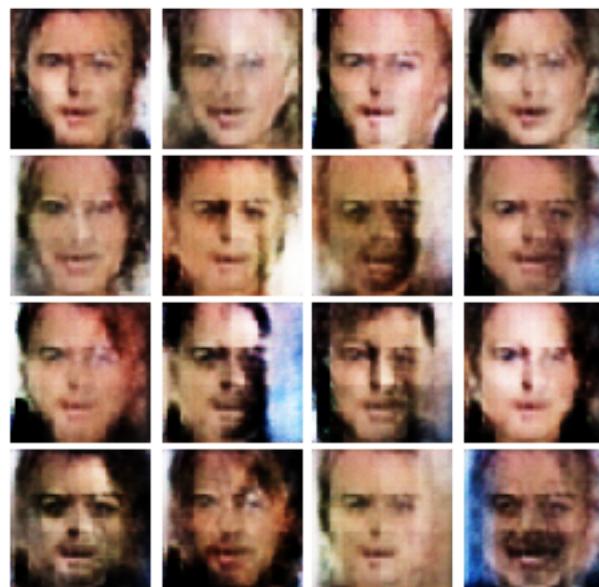


EPOCH: 2

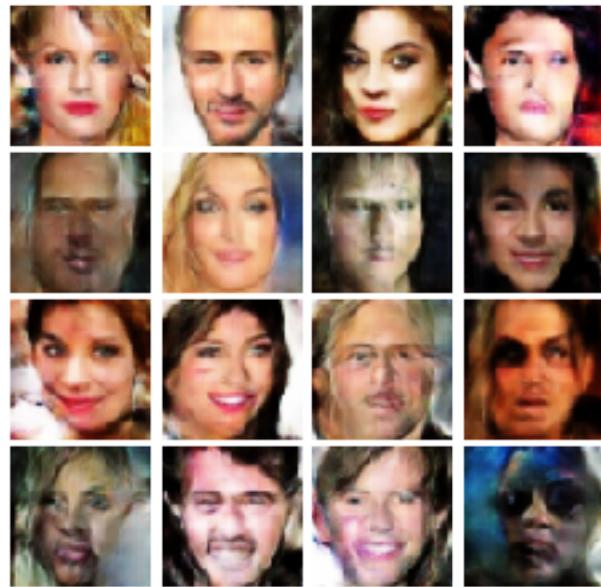
Iter: 1000, D: 0.4516, G:3.751



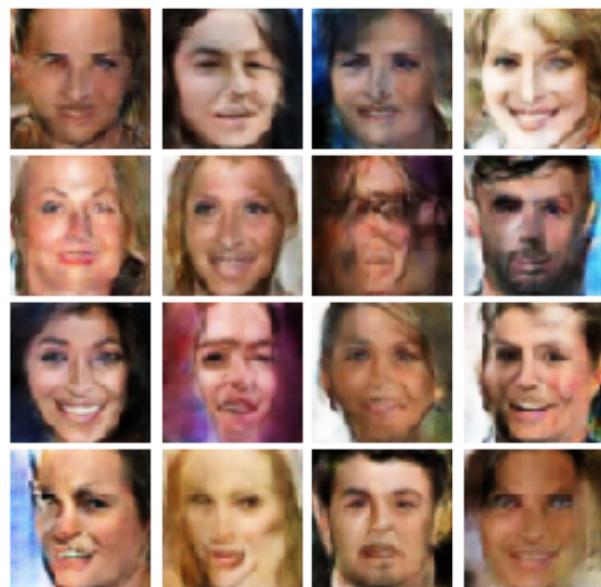
Iter: 1200, D: 1.25, G:6.047



Iter: 1400, D: 0.4507, G:2.75



Iter: 9600, D: 1.049, G:7.136



Iter: 9800, D: 0.08631, G:4.833



7 WGAN

```
[8]: from gan.wgan_models import w_Discriminator, w_Generator
      from gan.wgan_loss import w_discriminator_loss, w_generator_loss
      from gan.wgan_train import w_train
```

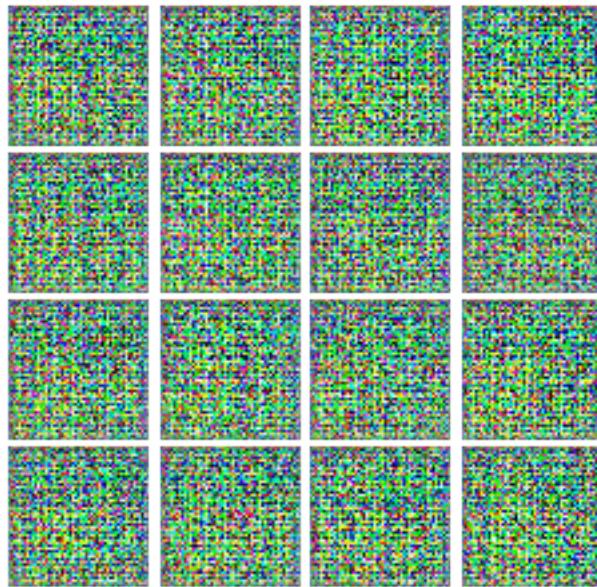
```
[15]: NOISE_DIM = 100
      NUM_EPOCHS = 20
      learning_rate = 0.0002
      c = 0.01
      n_critic = 2
```

```
[16]: D = w_Discriminator().to(device)
      G = w_Generator(noise_dim=NOISE_DIM).to(device)

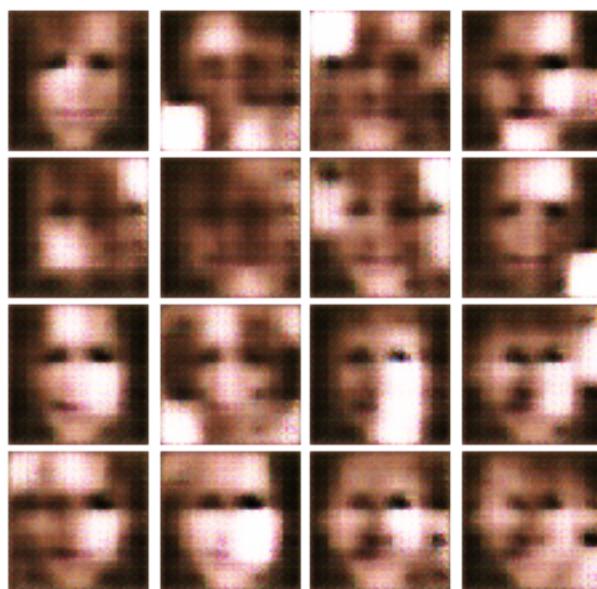
      D_optimizer = torch.optim.RMSprop(D.parameters(), lr=learning_rate)
      G_optimizer = torch.optim.RMSprop(G.parameters(), lr=learning_rate)

      w_train(D, G, D_optimizer, G_optimizer, w_discriminator_loss,
              w_generator_loss, c=c, num_epochs=NUM_EPOCHS, show_every=200,
              train_loader=celeba_loader_train, device=device)
```

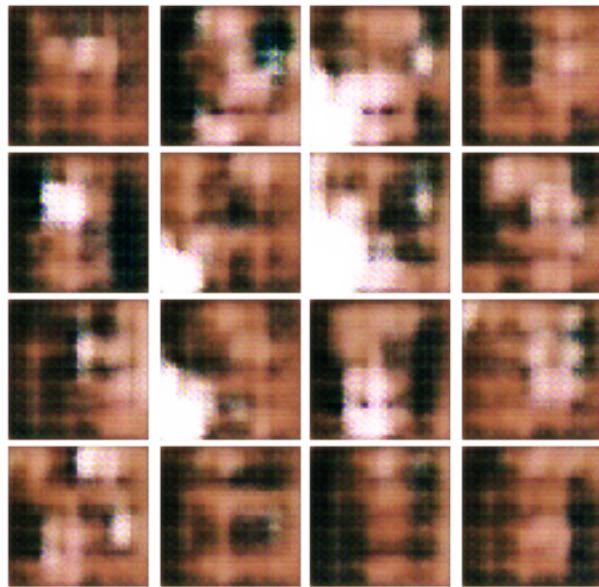
EPOCH: 1
Iter: 0, D: -0.05348, G: 0.1637



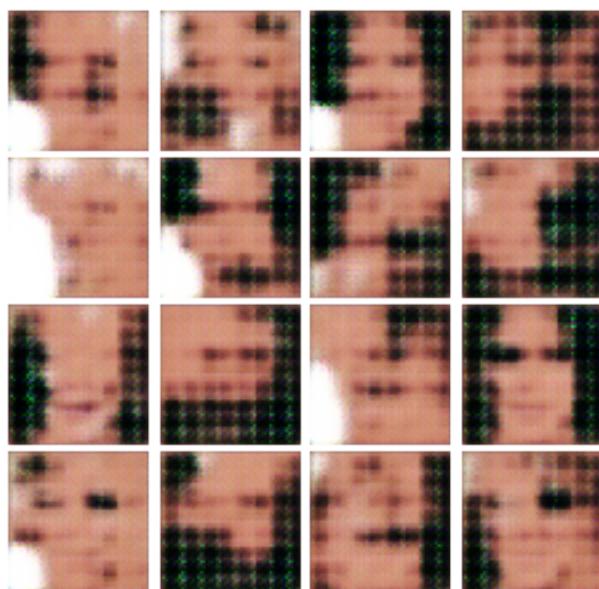
Iter: 200, D: -1.65, G:1.122



Iter: 400, D: -2.911, G:1.402



Iter: 600, D: -2.86, G:1.425

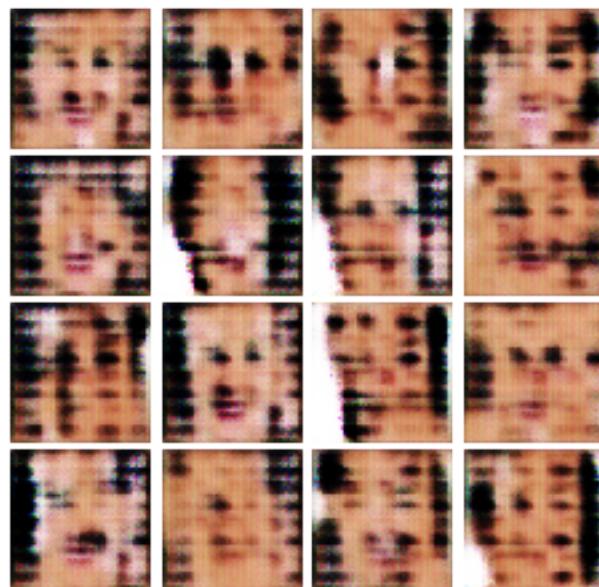


Iter: 800, D: -2.882, G:1.451

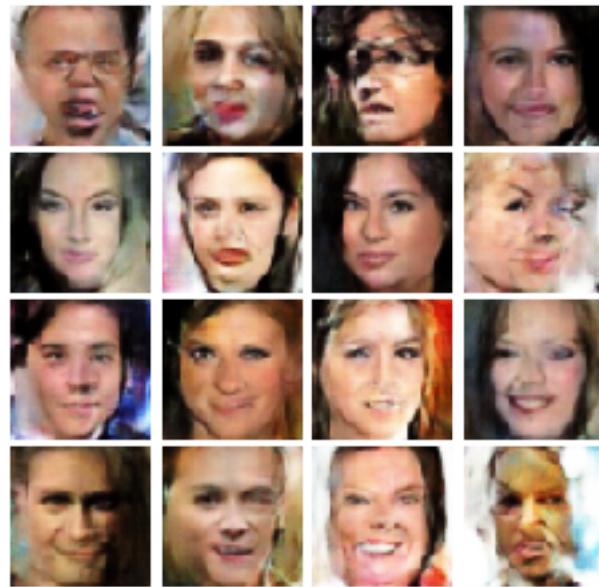


EPOCH: 2

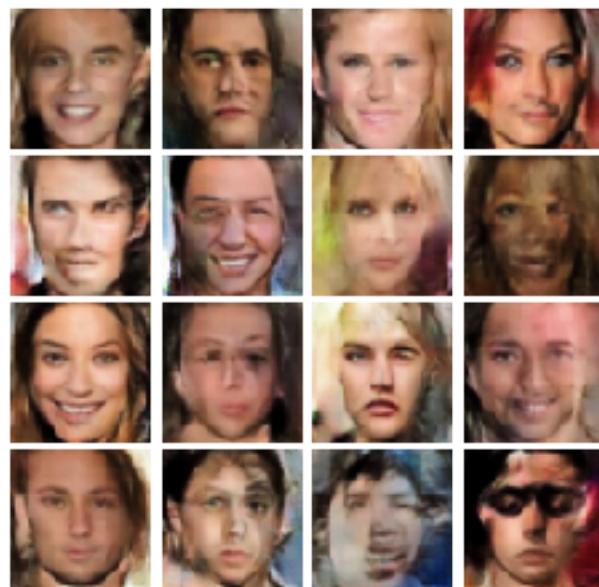
Iter: 1000, D: -2.951, G:1.448



Iter: 1200, D: -2.518, G:0.9893



Iter: 19400, D: -0.4049, G:0.8406



Iter: 19600, D: -0.5375, G:0.9865



[]: